

Database Design and Development

DATABASE DEVELOPMENT – SQL DML – JOINS

Specifying Joined Tables in the FROM Clause of SQL

Joined table

- It was incorporated into SQL to permit users to specify a table resulting from a join operation in the FROM clause of a query.
- JOIN may also be called INNER JOIN

fname	Lname	address
John	Smith	731 Fondren, Houston,TX
Franklin	Wong	638 Voss, Houston,TX
Joyce	English	5631 Rice, Houston, TX
Ramesh	Narayan	975 Fire Oak, Humble, TX

Join Example

Retrieve the name and address of all employees who work for the 'Research' department.

```

SELECT      employee.fname, employee.lname,
            employee.address
FROM        employee
JOIN        department
ON          department.number =
            employee.dept_number
WHERE       department.name = 'Research';

```

Multiway JOIN in the FROM clause

Can nest JOIN specifications for a multiway join:

For every project located in 'Stafford', list the project number, the controlling department number, and the department manager's last name, address, and birth date

```
SELECT      project.number, project.dept_number,  
              employee.lname, employee.address,  
              employee.birthdate  
  
FROM        project  
  
JOIN        department ON project.dept_number = department.number  
  
JOIN        employee ON department.manager_SSN = employee.SSN  
  
WHERE       project.location='Stafford';
```

Ambiguous Attribute Names

Same name can be used for two (or more) attributes in different relations

- As long as the attributes are in different relations
- Must **qualify** the attribute name with the relation's name to prevent ambiguity
- Examples:
 - dept_number attribute in project and employee relations
 - number and name attributes in project and department relations

For every project located in 'Stafford', list the project number, the controlling department number, and the department manager's last name, address, and birth date.

```
SELECT    project.number, project.dept_number,
            employee.lname, employee.address,
            employee.birthdate

FROM      project


JOIN      department

            ON project.dept_number = department.number

JOIN      employee

            ON department.manager_SSN = employee.SSN

WHERE     project.location = 'Stafford'
```



INNER and OUTER Joins

INNER JOIN (**versus** OUTER JOIN)

- Default type of join in a joined table
- Tuple is included in the result only if a matching tuple exists in the other relation


LEFT OUTER JOIN

- Every tuple in left table must appear in result
- If no matching tuple
 - Padded with NULL values for attributes of right table

RIGHT OUTER JOIN

- Every tuple in right table must appear in result
- If no matching tuple
 - Padded with NULL values for attributes of left table

Inner Join



Syntax:


```
SELECT      column_name(s)
FROM        table1 JOIN table2
ON table1.column_name = table2.column_name;
```

Example:

```
SELECT      *
FROM        employee JOIN department
ON employee.dept_number = department.number;
```

Note: JOIN is the same as INNER JOIN.

Left Outer Join



Syntax:

```
SELECT      column_name(s)
FROM        table1 LEFT JOIN table2
ON         table1.column_name = table2.column_name;
```


Example:

```
SELECT      *
FROM        employee LEFT JOIN department
ON         employee.dept_number = department.number;
```

Note: LEFT JOIN is the same as LEFT OUTER JOIN.

Right Outer Join

RIGHT JOIN



Syntax:

```
SELECT      column_name(s)
FROM        table1 RIGHT JOIN table2
ON          table1.column_name = table2.column_name;
```

Example:

```
SELECT      *
FROM        employee RIGHT JOIN department
ON          employee.SSN = department.manager_SSN;
```

Note: RIGHT JOIN is the same as RIGHT OUTER JOIN.

Aliasing / Tuple Variables

Aliases or tuple variables

- Declare alternative relation names E and S to refer to the EMPLOYEE relation twice in a query:
- Recommended practice to abbreviate names and to prefix same or similar attribute from multiple tables.

- For each employee, retrieve the employee's first and last name and his or her department name.

```
SELECT          e.fname, e.lname,  
                  d.name  
  
FROM            employee AS e  
  
JOIN            department AS d  
ON              e.dept_number =  
                  d.number;
```

Renaming of Attributes in SQL

Use qualifier AS followed by desired new name

- Rename any attribute that appears in the result of a query

For each employee, retrieve the employee's first and last name and the first and last name of his or her immediate supervisor.

```
SELECT      e.fname AS Employee_First_Name,  
              e.lname AS Employee_Last_Name,  
              s.fname AS Supervisor_First_Name,  
              s.lname AS Supervisor_Last_Name  
FROM        employee AS e  
JOIN        employee AS s  
ON          e.supervisor_SSN = s.SSN;
```

Arithmetic Operations

Standard arithmetic operators:

- Addition (+), subtraction (–), multiplication (*), and division (/) may be included as a part of **SELECT**

Show the resulting salaries if every employee working on the 'ProductX' project is given a 10 percent raise.

```
SELECT          e.fname, e.lname,  
                  1.1 * e.salary AS increased_sal  
FROM            employee AS e  
JOIN            employee_project AS w  
                  ON e.SSN=w.employee_SSN  
JOIN            project AS p  
                  ON w.project_number=p.number  
WHERE           p.name='ProductX';
```

Ordering of Query Results

Use **ORDER BY** clause

- Keyword **DESC** to see result in a descending order of values
- Keyword **ASC** to specify ascending order explicitly
- Typically placed at the end of the query

```
SELECT      *
FROM        employee e
JOIN        department d
ON          e.dept_number =
              d.number

ORDER BY    d.name ASC,
              e.salary DESC,
              e.fname ASC;
```

JOIN with Aggregate Functions

GROUP BY
may be
applied to
the result of
two relations

For each project, retrieve the project number, the project name, and the number of employees who work on that project.

```
SELECT      p.number, p.name, COUNT(*)
FROM        project p
JOIN        employee_project ep ON
           p.number = ep.project_number
GROUP BY    p.number, p.name;
```

JOIN with Aggregate Functions

Retrieve the total number of employees in the 'Research' department.

```
SELECT      COUNT (*)
FROM        employee
JOIN        department
ON dept_number = number
WHERE       name = 'Research';
```

Find the sum of the salaries of all employees of the 'Research' department, as well as the maximum salary, the minimum salary, and the average salary in this department.

```
SELECT  SUM(salary), MAX(salary),
          MIN(salary), AVG(salary)
FROM    employee
JOIN    department
ON dept_number = number
WHERE   name = 'Research';
```

HAVING Clauses

HAVING clause

- Provides a condition to select or reject an entire group:

For each project on which more than two employees work, retrieve the project number, the project name, and the number of employees who work on the project.

```
SELECT  p.number, p.name, COUNT(*)
FROM    project p
JOIN    employee_project ep
        ON p.number = ep.project_number
GROUP BY p.number, p.name
HAVING  COUNT(*) > 2;
```


References

Elmasri, R., & Navathe, S. (2017). Fundamentals of database systems (Vol. 7). Pearson

Nugent, D. (2017). Higher Nationals in Computing Core Textbook. Pearson Education Custom Content.